

Programming Assignment 2

CSE 4153/6153 – Data Communication Networks – Spring 2016

Building in Reliability

February 17, 2016

Due Date: Monday, March 14, 2016 by 11:59pm CST (local time in Starkville)

Assignments by undergraduate students are to be done in pairs or individually. No larger groups are allowed and you can only work on code with that particular partner. Graduate students must do the assignment individually.

Programs will be reviewed using the MOSS plagiarism-detection software.

Carefully read and follow the instructions for submitting your solution.

1. Assignment Objective

In the last assignment, you created a file-transfer protocol from client to server that was unreliable. The goal of this assignment is to incorporate some amount of reliability into your application.

To this end, you will implement a version of the **Go-Back-N (GBN)** protocol which facilitates the transfer of a text file from one host to another across an unreliable network. The protocol should be able to handle packet losses. For simplicity, we do not consider delayed packets; that is, a packet is received almost immediately after it is sent, or not at all.

Your protocol will remain unidirectional, i.e., data will flow in one direction (from the client to the server) and the acknowledgements (ACKs) in the opposite direction. To implement this protocol, you will write two programs: a client and a server, with the specifications given below. All communication will be done over datagram (UDP) sockets (as we saw in class).

You will test your implementation using an emulated network link (which will be provided to you as executable code) as shown in the diagram in Figure 1.

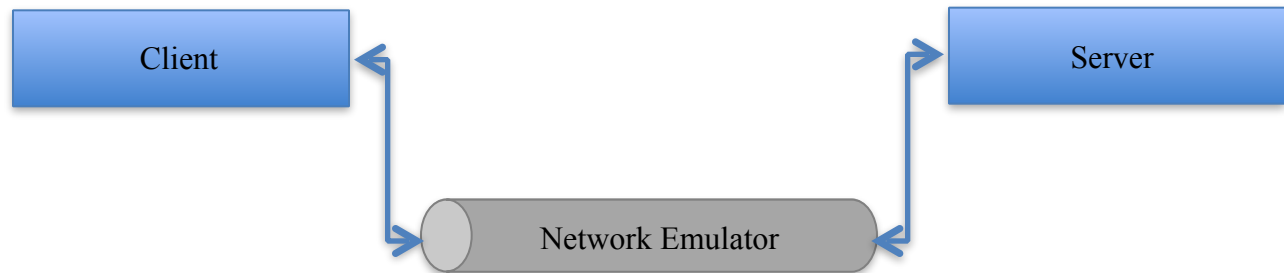


Figure 1

When the client sends packets to the server, it sends them to the network emulator instead of directly to the server. The network emulator then forwards the received packets to the server. However, it may randomly discard packets. The same scenario happens when the server sends ACKs to the client.

1.1 Using C++ versus Java

The assignment description (data structure and program names) sometimes differentiates between implementations in C++ and those in Java. Be sure to use the correct emulator and packet class depending on your choice of programming language. Below, I have summarized some the challenges facing you that are specific to C++ and those that are specific to Java.

C++: The emulator is an executable only; you do not see the code. Due to architectural issues, you will likely have to run the emulator on Pluto only. This means that you want your client and server code to be working before attaching them to the emulator (otherwise, you'll be fixing your client/server code locally and uploading to Pluto often). In Section 7, I provide advice on how to proceed. Conversely, for those comfortable developing code on Pluto, I recommend using the emacs editor.

On the plus (har har) side, The packet class already has the serialization and deserialization functions written and they are fairly straightforward.

Java: The emulator is an executable only, you do not see the code. However, given the great portability of Java (unlike C++), this should run locally as it does for me (java version "1.8.0_65"). Running on Pluto gives the irritating and familiar memory/thread error. For now, as you develop your code, make sure it compiles on Pluto even if it does not execute due to the memory error and you will not be penalized.

On the other hand, the serialization/deserialization commands require a little bit more effort to use. To smooth the way, I've provided you with links that should helpful.

1.2 Makefiles for C++ and Java

I am not providing makefiles this time. You can easily modify the one you used for PA1 to work for PA2, and you should submit your makefile with your source files (see Section 6).

2. Packet Format in Java and C++

A packet class is provided. All packets exchanged between the client and the server have the general structure as given by the code underneath. This class is similar for both C++ and Java, but I illustrate the Java class below (for both languages, be sure to understand the packet code before using it):

```
public class packet {  
    .....  
    private int type;  
    private int seqnum;  
    private int length;  
    private String data;  
    .....  
}
```

The `type` field indicates the type of the packet: it is set to 0 if it is an ACK, 1 if it is a data packet, 2 if it is an end-of-transmission (EOT) packet from server to client, and 3 if it is an EOT packet from client to server.

For data packets, `seqnum` is the modulo 8 sequence number of the packet. That is, sequence numbers have values in {0, 1, 2, 3, 4, 5, 6, 7}. The sequence number of the first packet should be zero.

For ACK packets, `seqnum` is the sequence number of the packet being acknowledged.

The `length` field specifies the number of characters carried in the data field. It should be in the range of 0 to 30, taking the largest chunk of data per packet sent (i.e. take 30 characters if you can, less if you have reached the end of the file). As with our previous assignment, you can assume you are working with an ASCII text file and that spaces, exclamation points, newline symbols, etc. each count as a character.

For ACK packets, `length` should be set to zero and these packets should carry no data (i.e. data field should be set to NULL).

You may **not** modify this class (we will use our own packet class when testing).

3. Client Program (`client`)

You should implement a client program, named `client`. Its command line input includes the following:

```
<emulatorName: host address of the emulator>,
```

<sendToEmulator: UDP port number used by the emulator to receive data from the client>,

<receiveFromEmulator: UDP port number used by the client to receive ACKs from the emulator>,

<fileName: name of the file to be transferred>

in the given order.

Upon execution, the client program should be able to read data from the specified file and send it using the GBN protocol to the server via the network emulator. The window size should be set to $N=7$. Recall that this means that, even though we have 8 sequence numbers, we are only using 7 of these sequence numbers at any given time.

After all contents of the file have been transmitted successfully to the server and corresponding ACKs have been received, the client should send an EOT packet to the server. The EOT packet is in the same format (and it has a sequence number) as a regular data packet, except that its type field is set to 3, its length is set to zero, and the data is set to NULL.

The client can close its connection and exit only after it has received ACKs for all data packets it has sent and received an EOT from the server. To keep the project simple, you can assume that EOT packets are never lost in the network.

To ensure reliable transmission, your program should implement the GBN protocol as follows:

- If the client has a packet to send, it first checks to see if the window is full — that is, whether there are N outstanding, unacknowledged packets. If the window is not full, the packet is sent and the appropriate variables are updated. If the window is full, the client will try sending the packet later.
- A timer of 800 milliseconds is started if it was not done before. The client will use only a single timer that will be set for the oldest transmitted-but-not-yet-acknowledged packet.
- When the client receives an acknowledgement packet with sequence number n , the ACK will be taken to be a cumulative acknowledgement, indicating that all packets with a sequence number up to and including n have been correctly received at the server.
- If a timeout occurs, the client resends all packets that have been previously sent but that have not yet been acknowledged. If an ACK is received corresponding to an un-acked packet within the window, but there are still additional transmitted-but-yet-to-be-acknowledged packets, the timer is restarted with 800 milliseconds. If there are no outstanding packets, the timer is stopped.

3.1 Output from Client

For grading purposes, your client program will generate two log files, named as *seqnum.log* and *ack.log*.

Whenever a packet is sent, its sequence number should be recorded in *seqnum.log*. This includes the EOT packet sent by the client.

The file *ack.log* should record the sequence numbers of all the ACK packets and the EOT packet (from the server) that the client receives during the entire period of transmission.

Each instance the program is executed, these files should be overwritten (not appended to).

The format for these two log files is one number per line. You must follow this format to avoid losing marks as the TA will be using these log files to grade the correctness of your assignment.

4. Server Program (server)

You should implement the server program, named `server`. Its command line input includes the following:

`<emulatorName: hostname for the emulator>,`

`<receiveFromEmulator: UDP port number used by the server to receive data from the emulator>,`

`<sendToEmulator-Port: UDP port number used by the emulator to receive ACKs from the server>,`

`<fileName: name of the file into which the received data is written>`

in the given order.

When receiving packets sent by the client via the network emulator, it should execute the following:

- check the sequence number of the packet.
- if the sequence number is the one that it is expecting, it should send an ACK packet back to the client with the sequence number of the received packet (indicating it has received all packets up to and including this packet).
- In all other cases, it should discard the received packet and resends an ACK packet for the most recently received in-order packet.

- After the server has received all data packets and an EOT from the client, it should send an EOT packet with the type field set to 2, and then exit.

Of course, the server must also write the received data from the client to the file `fileName`.

4.1 Output from Server

In addition to the file containing the text transferred from the client, the server program is also required to generate a log file, named as *arrival.log*. The file *arrival.log* should record the sequence numbers of all the data packets and the EOT packet (from the client) that the server receives during the entire period of transmission.

Each instance the program is executed, this file should be overwritten (not appended to). Similarly, in each instance the program is executed, `fileName` should be overwritten.

The format for the log file is one number per line. You must follow the format to avoid losing marks as the TA will be using these log files to grade the correctness of your assignment.

5. Network Emulator (emulator)

You will be given the **executable code** for the network emulator. To run it, you need to supply the following command line parameters in the given order:

<receivePort: emulator's receiving UDP port number, used by both client and server> ,

<sendToClient-Port: client's receiving UDP port number> ,

<sendToServer-Port: server's receiving UDP port number> ,

<clientName: client's network address> ,

<serverName: server's network address> ,

<seed: used to seed the random number generator> ,

<dropProb: packet discard probability> ,

<verbose-mode: int set to 1, the network emulator will output its internal processing>

When the network emulator receives a packet from the client, it will discard it with the specified probability. If not discarded, the packet is forwarded to the server.

Note that the emulator will not terminate. You do not need to do anything about this, it is acceptable behavior.

5.1 Students using Java

Serialization and deserialization should be used to transmit and receive packets, respectively. Learn about the following commands:

```
ByteArrayOutputStream oSt = new ByteArrayOutputStream();
ObjectOutputStream ooSt = new ObjectOutputStream(oSt);
ooSt.writeObject(pkt);
ooSt.flush();
byte[] sendBuf = new byte[30];
sendBuf = oSt.toByteArray();
```

and the commands for deserialization. The following link seems like a good start (or read about it anywhere else you would like):

<http://docs.oracle.com/javase/7/docs/api/java/io/ByteArrayOutputStream.html>

and you can see examples here:

<http://stackoverflow.com/questions/17940423/send-object-over-udp-in-java>

<http://stackoverflow.com/questions/3997459/send-and-receive-serialize-object-on-udp-in-java>

5.1 Students using C++

Serialization and deserialization functions have been provided to you in the packet class. Other methods of serializing/deserializing exist for C++, but they are complicated and they will not be used for this assignment.

5.2 How to Call Executables

For C++, you should call your executables in three separate terminals with the commands:

```
./emulator 6000 6001 6002 localhost localhost 0 0.0 1
./server localhost 6002 6000 output.txt
./client localhost 6000 6001 file.txt
```

Note that you should call the emulator and server first, then the client. Of course, you may change the parameter values for ports etc. — in fact, I suggest you choose random port values so your don't conflict with other students while testing.

For Java, you should call your executables similarly. For example:

```
java emulator 6000 6001 6002 localhost localhost 0 0.0 1
java server localhost 6002 6000 output.txt
java client localhost 6000 6001 file.txt
```

An example execution for both languages will be posted to MyCourses soon. My examples may produce additional output that is helpful. Your code need not do this; if you do decide to provide output to the screen/terminal, please be economical (do not spam).

6. Your Submitted Solution

6.1 Due Date

This assignment is due on Monday, March 14, 2016 by 11:59pm CST (local time in Starkville).

6.2 Hand in Instructions

Submit your files in a single compressed file (either .zip or .tar) using MyCourses.

Your code **MUST have the names of the author(s)** on it! If you are working as a pair, **only one student** should submit the code. You must hand in the following files:

- **Source code for client and server**
- **Your makefile: your code must compile and link cleanly by typing ``make'' (note that you DO NOT need to have anything regarding the emulator here)**

Your implementation will be tested on the Pluto system (with the Java exception noted earlier in Section 1.1) and so it must run smoothly in this environment! See the syllabus for penalties for late assignments.

A sample file will be provided online for you to use for testing your client/server programs. However, it is your responsibility to test your code thoroughly and ensure it conforms to the requirements.

6.3 Documentation

There is no external documentation required for this assignment, you are expected to have a reasonable amount of internal code documentation.

7. Suggestions for Getting Started

I recommend the following steps for implementation:

1. Forget about the emulator and GBN for now. Get the client and server programs running using serialization and deserializing such that they can exchange data using the two different ports assigned. Develop this locally, but remember to test that this works on Pluto.
2. Still ignoring GBN, get your client and server to exchange messages through the emulator with the drop probability set to 0. Break this task into two sub-pieces: (1) get the client to send data successfully to the emulator, (2) get the server to do the same. Note that you can modify your server code to send any message you like to the emulator for testing; the emulator will handle this and output its state if you set `vmode=1`.

If you are using C++, you will have to complete this on Pluto directly (see Section 1.1). Conversely, if you are using Java, you will have to complete this locally if you are receiving the thread error (see Section 1.1).

3. Review the details of GBN as given in the lecture slides and in our textbook (noting that we using the variation where the server sends an ACK after reception of each packet from the client).
4. Now forget the emulator and focus on getting your client and server to run the GBN protocol. Develop this locally, but remember to test this on Pluto.
5. Pull everything together and run GBN over your client and server using the emulator. If you are using C++, you will have to complete this on Pluto directly (see Section 1.1). Conversely, if you are using Java, you will have to complete this locally if you are receiving the thread error (see Section 1.1).

Remember to use port numbers at 1024 or above to avoid reserved port numbers.

Finally, get started early. You have a lot time to complete this assignment, but ***do not*** leave this to the last minute (or even the last few days). I will tend to ignore last-minute requests for help, especially in regard to debugging code, since there is more than an adequate amount of time to complete this assignment.

Acknowledgements: I am grateful to Dr. Raouf Boutaba who brought this assignment topic to my attention.